



Overall Risk Score **0.67** **HIGH**



Executive Summary

■■■■■■■■■■ ■■■■■■ declares 0 data type(s) collected and 2 shared in its Play Data Safety label, across 6 data type(s) analyzed. Our heterogeneous GNN audit model identifies **4 potential undeclared collection(s)** — data types implied by embedded SDK signals but absent from both label and policy — and **0 policy-vs-label mismatch(es)**. Overall risk score: **0.67 (HIGH)**.

Top discrepancies by risk:

- **installed_apps**: Undeclared Collection (confidence 1.00)
- **crash_logs**: Undeclared Collection (confidence 1.00)
- **diagnostics**: Undeclared Collection (confidence 1.00)

Class definitions:

CONSISTENT: Label, policy, and SDK signals agree.

POLICY_LABEL_MISMATCH: Label discloses data not mentioned in policy.

OVER_DISCLOSURE: Policy mentions data not declared in label.

UNDECLARED_COLLECTION: SDK signals collection undisclosed in label and policy.

Discrepancy Table

Sorted by risk class (undeclared collection first). Y/N: label=data-safety label, policy=policy text, sdk=SDK signal.

Data Type	Class	Conf.	Evidence
installed_apps	Undeclared Collection	1.00	label:N policy:N sdk:Y (1, 2)
crash_logs	Undeclared Collection	1.00	label:N policy:N sdk:Y (1, 2)
diagnostics	Undeclared Collection	1.00	label:N policy:N sdk:Y (1, 2)
device_id	Undeclared Collection	1.00	label:N policy:N sdk:Y (1, 2)
app_interactions	Consistent	0.98	label:N policy:N sdk:N (1, 2)
user_ids	Consistent	0.98	label:N policy:N sdk:N (1, 2)

Evidence Trails

Detailed evidence for UNDECLARED_COLLECTION or confidence > 0.85 non-CONSISTENT findings.

Discrepancy: installed_apps — Undeclared Collection (confidence 1.00)

Label declares collection: **No**

Label declares sharing: **No**

Policy text mentions: **No**

SDK collection signals: **Yes** — 1, 2

Recommended action: Add 'installed_apps' to the Play Data Safety label under 'Data Collected' and update the privacy policy to disclose this collection.

Discrepancy: crash_logs — Undeclared Collection (confidence 1.00)

Label declares collection: **No**

Label declares sharing: **No**

Policy text mentions: **No**

SDK collection signals: **Yes** — 1, 2

Recommended action: Add 'crash_logs' to the Play Data Safety label under 'Data Collected' and update the privacy policy to disclose this collection.

Discrepancy: diagnostics — Undeclared Collection (confidence 1.00)

Label declares collection: **No**

Label declares sharing: **No**

Policy text mentions: **No**

SDK collection signals: **Yes** — 1, 2

Recommended action: Add 'diagnostics' to the Play Data Safety label under 'Data Collected' and update the privacy policy to disclose this collection.

Discrepancy: device_id — Undeclared Collection (confidence 1.00)

Label declares collection: **No**

Label declares sharing: **No**

Policy text mentions: **No**

SDK collection signals: **Yes** — 1, 2

Recommended action: Add 'device_id' to the Play Data Safety label under 'Data Collected' and update the privacy policy to disclose this collection.

Methodology

PolicyGraphAudit-RT represents each Android app as a tri-partite heterogeneous graph connecting three layers: **Privacy Policy** text (OPP-115 segment classification), **Play Data Safety Labels** (declared data types and purposes), and **Runtime Evidence** (inferred SDK trackers via Exodus Privacy). A 2-layer HeteroConv R-GCN encodes all three layers jointly; an MLP classifier head predicts discrepancy classes for each (App, DataType) pair.

The model was trained under an edge-masking protocol (30% of label-determining edges removed) to prevent circularity. It achieves **macro F1 = 0.956** on 39 held-out apps (UNDECLARED_COLLECTION F1 = 0.974), trained on 252 apps and 3,202 labeled pairs. Model card: https://github.com/PolicyGraphAudit-RT/reports/blob/main/reports/m5_model_card.md.

Known limitations: (1) Privacy labels are from a 2022 Play Store snapshot. (2) SDK presence is inferred via category-level priors, not measured from APK bytecode. (3) Policy classifier uses MiniLM + logistic regression (OPP-115 macro F1 = 0.70). (4) Labels are rule-derived weak supervision, not human-annotated. (5) English-language Android apps only.

Disclaimer

This report is produced by a research prototype and is provided for informational and research purposes only. It does not constitute legal advice or a regulatory determination. Findings are based on automated weak-supervision labels from publicly available data and have not been reviewed by a privacy legal professional. Research prototype · Ongoing research.